Performance Impact of *CUDA* on Image Manipulation Algorithms

Luis Silva - 54449, Lourenço Soares - 54530

✦

1 INTRODUCTION

THIS report intends on investigating the impact of optimizations and implementations using CUDA on the processing time of image manipulation algorithms, which we will refer to as shaders (because that is, in essence, what the two algorithms are). This program in question implements two effects, the first being a blur operation (which requires reading neighboring pixels) and the second being a color correction (desaturation) shader.

We attempted to iterate on our solutions, experimenting with different approaches to setting up tasks and memory for the GPU in order to minimize the time spent performing the effects, and compared the results (which were obtained both with basic CPU timers and with a heavier tool like nvprof) with previous solutions.

2 APPROACH AND RESULTS

The programs used to benchmark different implementations, which are mentioned throughout this section, can be downloaded from our public *Git* repository[2]. The times given in this section were obtained via our autotester program, which tested each version of our implementations 10 times with the `tram.ppm` image on the University's cluster.

2.1 Original Code

First and foremost, after reading the original source code that was provided to us by the professor, we saw some improvements that could be made to the program in order to substantially reduce the time spent on computations. We started off by rewriting and commenting the source code to both have a better understanding of it, but also to allow for future improvements to be added in without much effort. We did this without actually modifying any of the logic to the original two shaders, because we did not want to alter the performance impact of the shaders themselves. Of note, our main changes were:

1) The blur shader was expanded to allow for larger blur operations to be performed (by modifying the `KERNELRADIUS` macro. This also required the blur matrix (confusingly named a Gaussian Kernel) to be generated dynamically from the given radius.
2) The blur matrix was packed into a struct to allow it to be properly passed by value, as opposed to passed by reference like C typically does for arrays.
3) Each pixel of a given image (called a *texel*) was put in a struct that contains an R, G, and B value, as opposed to just being an array which each element occupies 3 indices.

On average, this solution took 206.89ms to execute.

2.2 V01 - Basic CUDA Port

Next, we took the original source code and did a very basic port to CUDA with no changes to the underlying shader. This solution took, on average, 40.95ms to execute (9.56ms were spent just on executing the shaders themselves). Most of the time was being spent on copying memory to and from the GPU.

2.3 V02 - New Algorithms

Now having a benchmark to improve upon, now we could start improving the algorithms themselves. Our main point of contention with the original blur implementation was that we felt the `if` statements were unnecessary and could be avoided entirely (we had experience writing shaders, where avoiding branching is an important step in improving performance (after all, GPUs are designed for fast number crunching). We rewrote the logic to remove the if statements inside the loops themselves, opting instead to calculate the array boundaries before executing the loops. We also tried to use floats whenever possible, as integer and floating point numbers have similar performance on GPUs. By sticking to a single data type, less time needs to be spent on casting.

We made two other small improvements, the first was modifying the three division operations (used to normalize the color) into a single division and 3 multiplications. This was done because multiplication performance is significantly faster than division. The other minor improvement was in the desaturation algorithm, where we opted to pass `1-alpha` into the kernel instead of calculating each time (as that value was needed more often than just `alpha`).

While these improvements did increase the number of assembly instructions that are executed in the blur shader

by about 20, the amount of instructions per loop was decreased, resulting in a faster execution time of 37.98ms (5.63ms were spent executing the shaders).

2.4   V03 - Memory Alignment

The next obvious improvement that could be performed was modifying the image input and output data to reduce cache misses and ensure 32-bit alignment. The original image data used three sets of integers (each 4 bytes) per texel. We modified it instead to be a single byte per texel, with each array index occupying 4 bytes (the fourth byte was unused, merely existed to ensure alignment). The smaller data size would allow for more data to be packed per memory fetch, and the 32-bit alignment would avoid cases where a single texel's data would fall on two different cache lines.

Reducing the image data down to a single byte does mean that images will technically lose color depth, but considering that most image manipulation software only supports colors ranging from 0 to 255, we believed the trade-off was worth it.

This change resulted in our execution time being reduced to 15.07ms (3.45ms spent executing the shader). We were not surprised by the shader's time being reduced, but we were caught off guard by the entire execution time being reduced, until we remembered that the new data structure occupies significantly less memory and thus, the memory copy operations are much faster.

2.5   V04 - Combined Algorithms

The next obvious solution was to take both operations and to combine them into a single shader. This could be done since both shaders are associative, thus were not required to be done separately. Doing so would reduce overhead in the GPU having to set up the blocks and warps, as well as remove the need to recalculate texel indices.

This solution reduced our execution time down to 14.27ms (2.70ms spent executing the shader).

2.6   V05 - Multipass Blur

Having written blur shaders in the past, we remembered the fact that blur kernels are separable[1], meaning that we can reduce the blurring algorithm from $O(n^2)$ to $O(2n)$. This can be achieved by first performing a horizontal blur on the entire image, and then a vertical blur (or vice versa). This, unfortunately, undoes a bit of the work we did on the previous implementation, as the blur operation now requires two seperate shaders (however the second shader **also** performs the desaturation operation).

Our multipass blur resulted in an execution time of 14.13ms (2.86ms executing the shader). Because the times are so statistically small, we cannot say for certain whether this change had a big impact on performance compared to V04,

however it most definitely will have an impact if the size of KERNELRADIUS is increased.

For example, using KERNELRADIUS of size 10 (to perform a stronger blur) V04 took 105.88ms to execute in its entirety (93.64ms for the shader). V05, on the other hand, took 24.86ms with the same blur size (14.46ms for the shader). That was over four times faster.

2.7   V06 - Streams

The last of the significant improvements which we could do to the program was utilizing CUDA streams, which could theoretically improve our execution speed 3 times over. This is because memory copy operations could be interleaved with kernel executions. The trickiest part of implementing such a solution was the fact that we actually needed to copy more memory than just what was needed per texel, because of the fact that neighboring texels were needed as well. To ensure we were as efficient as possible, we wrote code to take this extra memory into account and perform the next memory copies at an offset.

Disappointingly, all this hard work only improved our overall execution time to 13.19ms (2.79ms for the shader).

2.7.1   V06 - Streams (Single Pass)

Out of curiosity, we wanted to know if using our single pass blur shader would result in improved execution time with streams. Surprisingly, it actually performed statistically worse, taking 14.47ms to execute in its entirety (2.87ms executing just the shader). We are not entirely sure as to why the increased execution time was so significant compared to the times between V04 and V05.

3   FURTHER DISCUSSION OF THE IMPLEMENTATIONS

During the presentation of our implementation (and the subsequent results), we glossed over a few design decisions, which we will cover in greater depth here.

3.1   Thread Block Sizes

In our final code, we settled on using the maximum available number of threads in the hardware, instead of having a static amount for every GPU. We assigned a single thread per texel, attempting to fill as many texels of the image into a single block.

```
1   threadCount = fmin(imgsize,
        properties.maxThreadsPerBlock);
2   blockCount = (int)ceil ((( float )(imgsize))/threadCount);
3   shader<<<blockCount, threadCount>>>(...);
```

We experimented with a few different block and grid sizes, including 2 dimensional grids. We either found no significant change in performance, or worse due to the increased overhead. With regards to 2 dimensional grids, while they

improved execution time of each kernel by around 10us due to the lack of needing to use division and modulo to calculate the x and y coordinate regarding which texel coordinate the thread would modify, it had a negative impact on the overall execution time due to the fact that it was very difficult to create 2 dimensional grids which would fit neatly in an arbitrarily sized image. Regardless of how we tried to shape it, our solution would always end up requiring more blocks than the 1 dimensional solution, with significantly more threads being left idle due to the coordinates ending up out of bounds.

3.2 Local and Global Shared Memory

Another possible improvement could be made by utilizing shared memory. We saw the potential for this, which could improve the performance of the blurring algorithm due to all the neighboring data it has to fetch. There was not any significant improvements that could be done to the color correction shader, as it would only need to access its own texels. We implemented a shared memory solution in a separate branch of V05, as implementing shared memory in the multipass blur shader would be significantly easier. Our solution was unfortunately unfinished, as the current version seems to be miscalculating the adjacent texel indices, resulting in an image that is correctly blurred, but has a few artefacts which we were unable to correct in time.

We tried both using local and global shared memory for this solution, but in either one, our execution time ended up being worse than the version that did not use shared memory by about 3ms in the execution of the shader. This was very likely due to the fact that the impact of calculating which neighboring texels to store in the shared array outweighed the performance improvements of not having to access global memory. It could be that our solutions were inefficient, and that with a bit more time this could be improved upon.

3.3 Stream Sizes

Our source code allows the number of streams to be changed easily by modifying the `STREAMCOUNT` macro, which is defaulted to `2`. We obtained this value by testing our program with a wide range of different stream sizes, and we plotted each size as a function of time. This yeilded the following graph:
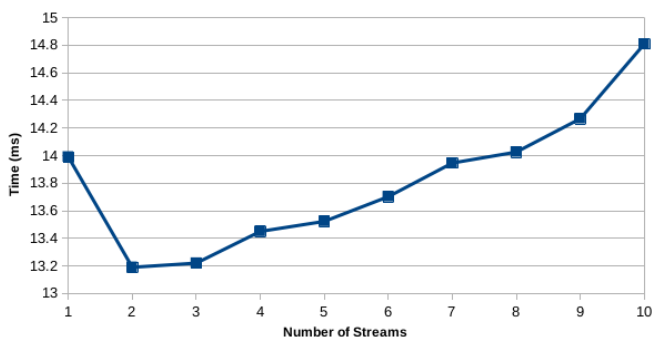

Fig. 1: The execution time of the program depending on the number of streams

A stream size of two seemed to yield the best overall execution time, and increasing the size further resulted in a worse performing program. This is likely due to the overhead involved in context switching outweighing the potential speedup increase. The execution time of the shaders remained relatively the same between all versions (averaging at around 2.84ms), which was expected as streaming mainly affects memory copying times.

3.4 Pinned Memory

While researching other solutions, we thought that looking into pinned memory could provide a possible decrease in execution time. Theoretically, it should have improved our execution times by preventing the need for paging. We modified our V06 solution to use `cudaMallocHost`, and this resulted in an average execution time of 68.80ms, with both shaders taking 30.378ms to run. We were surprised by this, as we were expecting better times, not significantly worse. It is possible we missed a caveat or did not implement it correctly, but we unfortunately did not have time to investigate this further.

4 CONCLUSION

After reviewing the results, we saw the significance that small changes to the program could bring. Simply porting a sequential program to CUDA already proved to be beneficial, and taking it steps further by being aware of how memory is stored and accessed, was a very educational experience. We can see the importance of parallel computing, especially in more real time applications such as real time rendering. We are relatively satisfied and confident with our results, having been able to bring the execution of the program down to 13.19ms, despite shortcomings due to time constraints.

5 FUTURE WORK

Other potential investigations we could have performed to improve our execution time would have been to reorganize our data to prevent memory bank conflicts. While this could be easily done for the color correction algorithm, this proved too difficult for the blurring algorithm as it requires neighboring texels. If we had more time, we could have potentially looked into a way to organize memory to allow for an easier time calculating the offsets inside the shader itself.

REFERENCES

[1] *Learn OpenGL - Bloom*. URL: https://learnopengl.com/Advanced-Lighting/Bloom.

[2] Luis Silva Lourenço Soares. *Project Repository*. URL: https://github.com/buu342/FCT-CAD_Project1.